

```
In [1]: import pandas as pd
```

```
In [2]: !pip install plotly
```

```
Requirement already satisfied: plotly in c:\users\ykyas\appdata\local\programs\python\python312\lib\site-packages (5.21.0)  
Requirement already satisfied: tenacity>=6.2.0 in c:\users\ykyas\appdata\local\programs\python\python312\lib\site-packages (from plotly) (8.2.3)  
Requirement already satisfied: packaging in c:\users\ykyas\appdata\local\programs\python\python312\lib\site-packages (from plotly) (24.0)
```

```
In [3]: import warnings
```

```
In [4]: warnings.filterwarnings('ignore')
```

```
In [5]: df=pd.read_csv('Titanic-Dataset.csv')
```

```
In [6]: df #we are printing the en
```

Out[6]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599 7
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803 5
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536 1
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053 3
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607 2
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369 3
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376

891 rows × 12 columns



In [7]:

```
df.info() #types of data
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age          714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

```

In [8]: `df.describe()` *#stats of the dataset*

Out[8]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch	
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.200000
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693000
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910000
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.450000
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.320000

In [9]: `df=df.drop(['PassengerId','Name','Ticket','Cabin'],axis=1)` *#retain whates needed*

In [10]: `df`

```
Out[10]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	22.0	1	0	7.2500	S
1	1	1	female	38.0	1	0	71.2833	C
2	1	3	female	26.0	0	0	7.9250	S
3	1	1	female	35.0	1	0	53.1000	S
4	0	3	male	35.0	0	0	8.0500	S
...
886	0	2	male	27.0	0	0	13.0000	S
887	1	1	female	19.0	0	0	30.0000	S
888	0	3	female	NaN	1	2	23.4500	S
889	1	1	male	26.0	0	0	30.0000	C
890	0	3	male	32.0	0	0	7.7500	Q

891 rows × 8 columns

```
In [11]: df.isna()
```

```
Out[11]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False
...
886	False	False	False	False	False	False	False	False
887	False	False	False	False	False	False	False	False
888	False	False	False	True	False	False	False	False
889	False	False	False	False	False	False	False	False
890	False	False	False	False	False	False	False	False

891 rows × 8 columns

```
In [12]: df.isna().sum() #to find the missing values
```

```
Out[12]: Survived      0
          Pclass       0
          Sex          0
          Age        177
          SibSp       0
          Parch       0
          Fare        0
          Embarked    2
          dtype: int64
```

```
In [13]: df=df.dropna(subset=['Embarked'])  #all missing embarked values are dropped
```

```
In [14]: df
```

```
Out[14]:
```


	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	22.0	1	0	7.2500	S
1	1	1	female	38.0	1	0	71.2833	C
2	1	3	female	26.0	0	0	7.9250	S
3	1	1	female	35.0	1	0	53.1000	S
4	0	3	male	35.0	0	0	8.0500	S
...
886	0	2	male	27.0	0	0	13.0000	S
887	1	1	female	19.0	0	0	30.0000	S
888	0	3	female	NaN	1	2	23.4500	S
889	1	1	male	26.0	0	0	30.0000	C
890	0	3	male	32.0	0	0	7.7500	Q

889 rows × 8 columns

```
In [15]: import seaborn as sss  #to use graphs
```


```
In [16]: sss.boxplot(df['Age'])  #outliers
```

```
Out[16]: <Axes: ylabel='Age'>
```




```
In [17]: sss.kdeplot(df['Age']) #normal
```

```
Out[17]: <Axes: xlabel='Age', ylabel='Density'>
```




```
In [18]: sss.boxplot(df['Fare']) #outliers
```

```
Out[18]: <Axes: ylabel='Fare'>
```




```
In [19]: sss.kdeplot(df['Fare']) #right sqewed
```

```
Out[19]: <Axes: xlabel='Fare', ylabel='Density'>
```




```
In [20]: sss.kdeplot(df['Pclass']) #normal graph
```

```
Out[20]: <Axes: xlabel='Pclass', ylabel='Density'>
```



```
In [21]: sss.boxplot(df['Pclass']) #no outliers
```

```
Out[21]: <Axes: ylabel='Pclass'>
```



```
In [22]: dfage=df['Age']  
dffare=df['Fare']  
#we need to fix age and fare
```



```
In [23]: import numpy as np
dfage_abs=dfage.abs() #avoid non negative numbers
dffare_abs=dffare.abs()

changedage=dfage_abs.apply(lambda x: x**(1/2)) #reducing the values using square
changedfare=dffare_abs.apply(lambda x: np.log1p(x))

df['Age']=changedage
df['Fare']=changedfare
```

```
In [24]: df['Age']
```


```
Out[24]: 0      4.690416
1      6.164414
2      5.099020
3      5.916080
4      5.916080
...
886    5.196152
887    4.358899
888         NaN
889    5.099020
890    5.656854
Name: Age, Length: 889, dtype: float64
```

```
In [25]: df['Fare']
```

```
Out[25]: 0      2.110213
1      4.280593
2      2.188856
3      3.990834
4      2.202765
...
886    2.639057
887    3.433987
888    3.196630
889    3.433987
890    2.169054
Name: Fare, Length: 889, dtype: float64
```


```
In [26]: sss.boxplot(df['Fare'])
```

```
Out[26]: <Axes: ylabel='Fare'>
```




```
In [27]: sss.kdeplot(df['Fare'])
```

```
Out[27]: <Axes: xlabel='Fare', ylabel='Density'>
```




```
In [28]: sss.boxplot(df['Age'])
```

```
Out[28]: <Axes: ylabel='Age'>
```



```
In [29]: sss.kdeplot(df['Age'])
```

```
Out[29]: <Axes: xlabel='Age', ylabel='Density'>
```



```
In [30]: df
```

Out[30]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	4.690416	1	0	2.110213	S
1	1	1	female	6.164414	1	0	4.280593	C
2	1	3	female	5.099020	0	0	2.188856	S
3	1	1	female	5.916080	1	0	3.990834	S
4	0	3	male	5.916080	0	0	2.202765	S
...
886	0	2	male	5.196152	0	0	2.639057	S
887	1	1	female	4.358899	0	0	3.433987	S
888	0	3	female	NaN	1	2	3.196630	S
889	1	1	male	5.099020	0	0	3.433987	C
890	0	3	male	5.656854	0	0	2.169054	Q

889 rows × 8 columns

In [31]: `from sklearn.preprocessing import OneHotEncoder` *#we are converting string values*

In [32]: `attr_val=['Sex','Embarked']
data_to_encode=df[attr_val]

data_to_encode=data_to_encode.dropna() #remove missing values as encoder will not work with them
encoder=OneHotEncoder(drop='first') #we remove sex_female, embarked_c because the first category is the baseline
encoded_data=encoder.fit_transform(data_to_encode) #transform data into encoded data
encoded_df=pd.DataFrame(encoded_data.toarray(), columns=encoder.get_feature_names_out())

df.reset_index(drop=True, inplace=True) #reset indices
encoded_df.reset_index(drop=True, inplace=True)

df_encoded=pd.concat([df, encoded_df], axis=1) #joining old dataframe with new data
df=df_encoded.drop(columns=attr_val) #remove original columns that got encoded`

In [33]: `df`

```
Out[33]:
```


	Survived	Pclass	Age	SibSp	Parch	Fare	Sex_male	Embarked_Q	Embarked
0	0	3	4.690416	1	0	2.110213	1.0	0.0	
1	1	1	6.164414	1	0	4.280593	0.0	0.0	
2	1	3	5.099020	0	0	2.188856	0.0	0.0	
3	1	1	5.916080	1	0	3.990834	0.0	0.0	
4	0	3	5.916080	0	0	2.202765	1.0	0.0	
...
884	0	2	5.196152	0	0	2.639057	1.0	0.0	
885	1	1	4.358899	0	0	3.433987	0.0	0.0	
886	0	3	NaN	1	2	3.196630	0.0	0.0	
887	1	1	5.099020	0	0	3.433987	1.0	0.0	
888	0	3	5.656854	0	0	2.169054	1.0	1.0	

889 rows × 9 columns




```
In [34]: sss.boxplot(df['Age'])
```

```
Out[34]: <Axes: ylabel='Age'>
```



```
In [35]: sss.kdeplot(df['Age'])
```

```
Out[35]: <Axes: xlabel='Age', ylabel='Density'>
```



```
In [36]: df.isna().sum() #we still have outliers in age
```

```
Out[36]: Survived      0
Pclass      0
Age         177
SibSp       0
Parch       0
Fare        0
Sex_male    0
Embarked_Q  0
Embarked_S  0
dtype: int64
```

```
In [37]: #we could not eliminate missing values,so we use linear regression to predict mi
from sklearn.linear_model import LinearRegression
from sklearn.impute import SimpleImputer
```

```
In [38]: #seperating rows with values and without values
without_age=df[df['Age'].isnull()]
with_age=df.dropna(subset=['Age'])

#training using with_age
x_train=with_age.drop(columns=['Age'])
y_train=with_age['Age']
regressor_obj=LinearRegression()
regressor_obj.fit(x_train,y_train)

#predicting missing data
missing_x_values=without_age.drop(columns=['Age'])
predicted_age=regressor_obj.predict(missing_x_values)

#replacing missing values with predicted values in the dataset
df.loc[df['Age'].isnull(),'Age']=predicted_age
```

In [39]: df

```
Out[39]:
```

	Survived	Pclass	Age	SibSp	Parch	Fare	Sex_male	Embarked_Q	Embarked_S
0	0	3	4.690416	1	0	2.110213	1.0	0.0	
1	1	1	6.164414	1	0	4.280593	0.0	0.0	
2	1	3	5.099020	0	0	2.188856	0.0	0.0	
3	1	1	5.916080	1	0	3.990834	0.0	0.0	
4	0	3	5.916080	0	0	2.202765	1.0	0.0	
...
884	0	2	5.196152	0	0	2.639057	1.0	0.0	
885	1	1	4.358899	0	0	3.433987	0.0	0.0	
886	0	3	4.583992	1	2	3.196630	0.0	0.0	
887	1	1	5.099020	0	0	3.433987	1.0	0.0	
888	0	3	5.656854	0	0	2.169054	1.0	1.0	

889 rows × 9 columns



```
In [40]: df.isna().sum() #no outliers left in age
```

```
Out[40]: Survived      0
Pclass      0
Age          0
SibSp       0
Parch       0
Fare        0
Sex_male    0
Embarked_Q  0
Embarked_S  0
dtype: int64
```

```
In [41]: #remove any duplicates if present
df=df.drop_duplicates()
```

```
In [42]: df #around 100 rows were removed
```

```
Out[42]:
```

	Survived	Pclass	Age	SibSp	Parch	Fare	Sex_male	Embarked_Q	Embarked
0	0	3	4.690416	1	0	2.110213	1.0	0.0	
1	1	1	6.164414	1	0	4.280593	0.0	0.0	
2	1	3	5.099020	0	0	2.188856	0.0	0.0	
3	1	1	5.916080	1	0	3.990834	0.0	0.0	
4	0	3	5.916080	0	0	2.202765	1.0	0.0	
...
883	0	3	6.244998	0	5	3.405355	0.0	1.0	
885	1	1	4.358899	0	0	3.433987	0.0	0.0	
886	0	3	4.583992	1	2	3.196630	0.0	0.0	
887	1	1	5.099020	0	0	3.433987	1.0	0.0	
888	0	3	5.656854	0	0	2.169054	1.0	1.0	

778 rows × 9 columns




```
In [43]: #we have finished the preprocessing of data ,now we need to do visualization
```

```
In [44]: import matplotlib.pyplot as plt
import seaborn as sss
```

```
In [45]: sss.countplot(df,x='Pclass',hue='Survived')
```

```
Out[45]: <Axes: xlabel='Pclass', ylabel='count'>
```

In [46]: *#we can see 3rd class people have not survived ,so 3rd class is in lower berth o
#1st and second class have survived*

df

Out[46]:

	Survived	Pclass	Age	SibSp	Parch	Fare	Sex_male	Embarked_Q	Emba
--	----------	--------	-----	-------	-------	------	----------	------------	------

0	0	3	4.690416	1	0	2.110213	1.0	0.0
1	1	1	6.164414	1	0	4.280593	0.0	0.0
2	1	3	5.099020	0	0	2.188856	0.0	0.0
3	1	1	5.916080	1	0	3.990834	0.0	0.0
4	0	3	5.916080	0	0	2.202765	1.0	0.0
...
883	0	3	6.244998	0	5	3.405355	0.0	1.0
885	1	1	4.358899	0	0	3.433987	0.0	0.0
886	0	3	4.583992	1	2	3.196630	0.0	0.0
887	1	1	5.099020	0	0	3.433987	1.0	0.0
888	0	3	5.656854	0	0	2.169054	1.0	1.0

778 rows × 9 columns



In [49]: `df['Family']=df['SibSp']+df['Parch']`
`df['Family']` *#we created dataset of people containing family*


```
Out[49]: 0      1
         1      1
         2      0
         3      1
         4      0
         ..
        883     5
        885     0
        886     3
        887     0
        888     0
        Name: Family, Length: 778, dtype: int64
```

```
In [50]: df['Alone']=df['Family']==0 #create dataset of people who dont have family
```

```
In [51]: df=df.drop(['SibSp','Parch'],axis=1)
```

```
In [52]: sss.boxplot(df)
```

```
Out[52]: <Axes: >
```




```
In [53]: #outliers in family ,age,embarked,fare
#standardize the age first

from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
scale_age=['Age']

df[scale_age]=scaler.fit_transform(df[scale_age])
```

```
In [54]: sss.boxplot(df)
```

```
Out[54]: <Axes: >
```



```
In [55]: rel_mat=df.corr()
rel_with_surv=rel_mat['Survived'].drop('Survived')
```

```
In [56]: rel_with_surv
```

```
Out[56]: Pclass      -0.333291
Age         -0.115660
Fare         0.309705
Sex_male    -0.511686
Embarked_Q  -0.038347
Embarked_S  -0.131168
Family       0.015938
Alone       -0.179048
Name: Survived, dtype: float64
```

```
In [57]: #lets assign a threshold and find values nearest to target variable
threshold=0.2
close_rel_with_surv=rel_with_surv[abs(rel_with_surv)>threshold]
```

```
In [59]: close_rel_with_surv
```

```
Out[59]: Pclass      -0.333291
Fare         0.309705
Sex_male    -0.511686
Name: Survived, dtype: float64
```

```
In [60]: ID=df[['Pclass','Age','Fare','Sex_male']]
D=df['Survived']
```

```
In [61]: ID
```

Out[61]:

	Pclass	Age	Fare	Sex_male
0	3	-0.383212	2.110213	1.0
1	1	0.638601	4.280593	0.0
2	3	-0.099957	2.188856	0.0
3	1	0.466450	3.990834	0.0
4	3	0.466450	2.202765	1.0
...
883	3	0.694464	3.405355	0.0
885	1	-0.613028	3.433987	0.0
886	3	-0.456988	3.196630	0.0
887	1	-0.099957	3.433987	1.0
888	3	0.286748	2.169054	1.0

778 rows × 4 columns

In [62]:

D

Out[62]:

```
0      0
1      1
2      1
3      1
4      0
..
883    0
885    1
886    0
887    1
888    0
```

Name: Survived, Length: 778, dtype: int64

In [63]:

```
from sklearn.model_selection import train_test_split
ID_train, ID_test, D_train, D_test = train_test_split(ID, D, test_size=0.2, random_state
```

In [64]:

ID_train

Out[64]:

	Pclass	Age	Fare	Sex_male
788	1	1.066956	4.384524	1.0
722	2	1.267117	2.639057	1.0
141	3	-0.238630	2.824351	0.0
388	2	-0.776486	2.564949	0.0
56	2	-0.457969	2.442347	0.0
...
72	3	-0.099957	2.737881	1.0
112	3	-0.534529	2.381858	0.0
287	2	0.857887	2.639057	1.0
483	1	-0.168600	4.522649	1.0
108	3	-0.637370	3.224858	0.0

622 rows × 4 columns

In [65]: ID_test

Out[65]:

	Pclass	Age	Fare	Sex_male
676	3	-0.693621	2.383400	0.0
667	3	0.911056	2.202765	1.0
614	2	-0.238630	4.189655	0.0
728	3	-0.168600	2.188856	0.0
545	2	-0.613028	3.295837	0.0
...
382	1	0.466450	3.970292	0.0
82	1	0.033477	3.873282	1.0
156	3	0.162225	2.202765	1.0
362	3	0.466450	2.085672	1.0
177	2	0.162225	2.639057	1.0

156 rows × 4 columns

In [66]: D_train

```
Out[66]: 788    0
          722    0
          141    1
          388    1
           56    1
          ..
          72     0
          112    0
          287    1
          483    1
          108    1
          Name: Survived, Length: 622, dtype: int64
```

```
In [67]: D_test
```

```
Out[67]: 676    1
          667    0
          614    1
          728    0
          545    1
          ..
          382    1
           82     0
          156    0
          362    0
          177    0
          Name: Survived, Length: 156, dtype: int64
```

```
In [80]: #using logistic regression model for training
         #first technique
         from sklearn.linear_model import LogisticRegression
         from sklearn.metrics import accuracy_score, confusion_matrix
         from sklearn.model_selection import cross_val_score
         from sklearn.metrics import classification_report

         #training the model
         model1=LogisticRegression()
         model1.fit(ID_train,D_train)
         D_pred=model1.predict(ID_test)

         #lets determine the model accuracy
         accuracy=accuracy_score(D_test,D_pred)
         print(accuracy)
```

```
0.782051282051282
```

```
In [82]: #using decsion tree second method
         from sklearn.tree import DecisionTreeClassifier
         from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
         from sklearn.model_selection import cross_val_score

         model2=DecisionTreeClassifier()
         model2.fit(ID_train,D_train)
         D_pred=model2.predict(ID_test)

         accuracy=accuracy_score(D_test,D_pred)
         print(accuracy)
```

```
0.7243589743589743
```

```
In [85]: !pip install xgboost
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matri
from sklearn.model_selection import cross_val_score

model3=XGBClassifier()
model3.fit(ID_train,D_train)
D_pred=model3.predict(ID_test)

accuracy=accuracy_score(D_test,D_pred)
print(accuracy)
```

Collecting xgboost

Downloading xgboost-2.0.3-py3-none-win_amd64.whl.metadata (2.0 kB)

Requirement already satisfied: numpy in c:\users\ykyas\AppData\Local\Programs\Python\Python312\lib\site-packages (from xgboost) (1.26.4)

Requirement already satisfied: scipy in c:\users\ykyas\AppData\Local\Programs\Python\Python312\lib\site-packages (from xgboost) (1.13.0)

Downloading xgboost-2.0.3-py3-none-win_amd64.whl (99.8 MB)

```
----- 0.0/99.8 MB ? eta -:--:--
----- 0.0/99.8 MB 660.6 kB/s eta 0:02:31
----- 0.4/99.8 MB 5.5 MB/s eta 0:00:19
----- 1.4/99.8 MB 11.1 MB/s eta 0:00:09
----- 2.3/99.8 MB 13.2 MB/s eta 0:00:08
----- 3.5/99.8 MB 15.8 MB/s eta 0:00:07
----- 4.6/99.8 MB 17.5 MB/s eta 0:00:06
----- 5.8/99.8 MB 18.6 MB/s eta 0:00:06
----- 6.8/99.8 MB 19.7 MB/s eta 0:00:05
----- 8.1/99.8 MB 19.8 MB/s eta 0:00:05
----- 9.2/99.8 MB 20.2 MB/s eta 0:00:05
----- 10.4/99.8 MB 24.2 MB/s eta 0:00:04
----- 10.6/99.8 MB 21.8 MB/s eta 0:00:05
----- 11.8/99.8 MB 23.4 MB/s eta 0:00:04
----- 13.3/99.8 MB 23.4 MB/s eta 0:00:04
----- 14.4/99.8 MB 24.2 MB/s eta 0:00:04
----- 15.7/99.8 MB 24.3 MB/s eta 0:00:04
----- 17.0/99.8 MB 25.2 MB/s eta 0:00:04
----- 18.4/99.8 MB 24.2 MB/s eta 0:00:04
----- 19.1/99.8 MB 25.2 MB/s eta 0:00:04
----- 20.3/99.8 MB 24.2 MB/s eta 0:00:04
----- 21.3/99.8 MB 26.2 MB/s eta 0:00:03
----- 22.3/99.8 MB 25.1 MB/s eta 0:00:04
----- 23.4/99.8 MB 24.2 MB/s eta 0:00:04
----- 24.8/99.8 MB 25.2 MB/s eta 0:00:03
----- 26.1/99.8 MB 25.2 MB/s eta 0:00:03
----- 27.1/99.8 MB 24.2 MB/s eta 0:00:03
----- 28.1/99.8 MB 24.2 MB/s eta 0:00:03
----- 29.1/99.8 MB 23.4 MB/s eta 0:00:04
----- 30.3/99.8 MB 23.4 MB/s eta 0:00:03
----- 31.5/99.8 MB 24.2 MB/s eta 0:00:03
----- 32.9/99.8 MB 25.1 MB/s eta 0:00:03
----- 33.8/99.8 MB 25.2 MB/s eta 0:00:03
----- 35.0/99.8 MB 24.2 MB/s eta 0:00:03
----- 36.1/99.8 MB 25.1 MB/s eta 0:00:03
----- 37.5/99.8 MB 25.2 MB/s eta 0:00:03
----- 38.9/99.8 MB 25.2 MB/s eta 0:00:03
----- 40.0/99.8 MB 26.2 MB/s eta 0:00:03
----- 41.5/99.8 MB 26.2 MB/s eta 0:00:03
----- 42.8/99.8 MB 27.3 MB/s eta 0:00:03
----- 43.8/99.8 MB 27.3 MB/s eta 0:00:03
----- 45.0/99.8 MB 28.4 MB/s eta 0:00:02
----- 46.0/99.8 MB 26.2 MB/s eta 0:00:03
----- 47.0/99.8 MB 27.3 MB/s eta 0:00:02
----- 48.0/99.8 MB 26.2 MB/s eta 0:00:02
----- 49.0/99.8 MB 25.2 MB/s eta 0:00:03
----- 50.4/99.8 MB 25.2 MB/s eta 0:00:02
----- 51.7/99.8 MB 25.2 MB/s eta 0:00:02
----- 52.9/99.8 MB 24.2 MB/s eta 0:00:02
----- 54.4/99.8 MB 26.2 MB/s eta 0:00:02
----- 55.7/99.8 MB 27.3 MB/s eta 0:00:02
----- 57.2/99.8 MB 27.3 MB/s eta 0:00:02
----- 58.5/99.8 MB 28.5 MB/s eta 0:00:02
```



```

----- 59.9/99.8 MB 29.7 MB/s eta 0:00:02
----- 61.1/99.8 MB 29.7 MB/s eta 0:00:02
----- 62.6/99.8 MB 29.7 MB/s eta 0:00:02
----- 63.9/99.8 MB 29.7 MB/s eta 0:00:02
----- 65.1/99.8 MB 28.4 MB/s eta 0:00:02
----- 66.5/99.8 MB 28.4 MB/s eta 0:00:02
----- 67.5/99.8 MB 27.3 MB/s eta 0:00:02
----- 68.4/99.8 MB 26.2 MB/s eta 0:00:02
----- 69.9/99.8 MB 27.3 MB/s eta 0:00:02
----- 70.8/99.8 MB 26.2 MB/s eta 0:00:02
----- 71.9/99.8 MB 26.2 MB/s eta 0:00:02
----- 73.3/99.8 MB 26.2 MB/s eta 0:00:02
----- 74.6/99.8 MB 26.2 MB/s eta 0:00:01
----- 75.6/99.8 MB 26.2 MB/s eta 0:00:01
----- 76.8/99.8 MB 24.3 MB/s eta 0:00:01
----- 77.7/99.8 MB 25.1 MB/s eta 0:00:01
----- 78.9/99.8 MB 24.2 MB/s eta 0:00:01
----- 80.3/99.8 MB 26.2 MB/s eta 0:00:01
----- 81.6/99.8 MB 26.2 MB/s eta 0:00:01
----- 82.6/99.8 MB 26.2 MB/s eta 0:00:01
----- 83.8/99.8 MB 26.2 MB/s eta 0:00:01
----- 84.4/99.8 MB 24.2 MB/s eta 0:00:01
----- 85.8/99.8 MB 24.2 MB/s eta 0:00:01
----- 86.9/99.8 MB 24.2 MB/s eta 0:00:01
----- 88.0/99.8 MB 25.1 MB/s eta 0:00:01
----- 89.2/99.8 MB 25.2 MB/s eta 0:00:01
----- 90.2/99.8 MB 24.2 MB/s eta 0:00:01
----- 91.8/99.8 MB 24.2 MB/s eta 0:00:01
----- 92.9/99.8 MB 25.2 MB/s eta 0:00:01
----- 94.0/99.8 MB 24.2 MB/s eta 0:00:01
----- 95.3/99.8 MB 25.2 MB/s eta 0:00:01
----- 96.6/99.8 MB 26.2 MB/s eta 0:00:01
----- 97.3/99.8 MB 26.2 MB/s eta 0:00:01
----- 97.8/99.8 MB 24.2 MB/s eta 0:00:01
----- 99.3/99.8 MB 25.1 MB/s eta 0:00:01
----- 99.7/99.8 MB 24.2 MB/s eta 0:00:01
----- 99.7/99.8 MB 24.2 MB/s eta 0:00:01
----- 99.7/99.8 MB 24.2 MB/s eta 0:00:01
----- 99.7/99.8 MB 24.2 MB/s eta 0:00:01
----- 99.7/99.8 MB 24.2 MB/s eta 0:00:01
----- 99.7/99.8 MB 24.2 MB/s eta 0:00:01
----- 99.8/99.8 MB 13.3 MB/s eta 0:00:00

```

Installing collected packages: xgboost
 Successfully installed xgboost-2.0.3
 0.7692307692307693

```

In [93]: #using random forest method
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.model_selection import cross_val_score

model4=RandomForestClassifier()
model4.fit(ID_train,D_train)

D_pred=model4.predict(ID_test)

accuracy=accuracy_score(D_test,D_pred)
print(accuracy)

```

```
print(classification_report(D_test,D_pred))
```

0.782051282051282

	precision	recall	f1-score	support
0	0.78	0.87	0.82	91
1	0.78	0.66	0.72	65
accuracy			0.78	156
macro avg	0.78	0.76	0.77	156
weighted avg	0.78	0.78	0.78	156

In [94]: *#at the moment we have random forest giving best accuracy*

In []: